

# Intro to CUDA/CUPY

# GPUs are Fast

- Graphics cards have to pump a lot of pixels, with little pixels (mostly) independent. Usually doing basic matrix ops to render.
- Same hardware turned out to be great for massively parallel numerical computations.
- NVidia's CUDA was phase change - made power of GPUs accessible for general purpose programming (GPGPU). See also OpenCL, others...
- Modern GPUs have order 10k cores, ~1 TB/s memory bandwidth, 50 TFLOPS (single precision), ~1 PFLOP (tensor cores)
- Memory bandwidth large: ~1 TB/s. Latency very large (often many hundreds of clock cycles).
- Double precision often much, much slower than single. Pay close attention to datatypes.

# If You Write Yourself

- You must understand how to split up your job into (tens of) thousands of streams.
- You can, and probably must, handle memory yourself.
- Threads come in blocks, and some resources shared within block. You can manually manage shared memory yourself.
- Reasoning about thread behavior when sharing memory can get very tricky.

# Fortunately...

- Many basic libraries ported to CUDA. NVidia supports linear algebra, FFTs. As always, steal if possible.
- High level languages offer GPU support. e.g. cupy, Jax work in python, see also tensorflow, pytorch.
- We'll stick with cupy - very close to drop-in replacement for numpy. Often same code will work for both.

# Memory Transfers

- CPU can't directly see memory on GPU. You have to transfer data to/from GPU.
- Bandwidth between CPU, GPU much slower than internal bandwidth for both.
- If it takes longer to send your data to GPU than to do problem on CPU, no point in using GPU.

# Profiling

- In general, GPU calls are asynchronous. You tell GPU to do something, it says “OK, I’ll go do that”, and returns immediately.
- This is a good thing, but if you aren’t careful you may be confused about run time.
- `cupy` comes with built-in profiler (`cupyx.profiler.benchmark`). Suggest you use it.

# Matrix Benchmark Revisited

- We'll time our matrix benchmark using cupy.
- Many numpy calls are directly supported. Just replace numpy with cupy
- Cupy allocated directly on gpu. If you want to transfer from gpu, `cupy.asnumpy(cupy_array)` will transfer to numpy array on CPU.
- `cupy.asarray(numpy_array)` will transfer to GPU



```

from cupy import benchmark
import cupy as cp

n=[100,300,1000,3000,5000,7000,10000,15000]
nrep=[1000,200,20,3,3,1,1,1]

for i,nn in enumerate(n):
    x=cp.random.rand(nn,nn,dtype='float32')
    #x=cp.random.rand(nn,nn) #changed one character!
    t1=time.time()
    for j in range(nrep[i]):
        y=x@x
    print(y[0,0])
    t2=time.time()
    nop=2*nn**3*nrep[i]
    gflops=nop/(t2-t1)/1e9
    print('got ',gflops,' gigaflops on matrix size of ',nn,
    print(benchmark(cp.matmul,(x,x),n_repeat=nrep[i]))

```

```

>>> exec(open("time_matrix_mult_gpu.py").read())
got 50.739185135004355 gigaflops on matrix size of 100
got 269.8426353726098 gigaflops on matrix size of 300
got 484.7687431195046 gigaflops on matrix size of 1000
got 543.5530987963836 gigaflops on matrix size of 3000
got 533.7326663982502 gigaflops on matrix size of 5000
got 548.506438242777 gigaflops on matrix size of 7000
got 551.3213171786251 gigaflops on matrix size of 10000
got 554.5790350019305 gigaflops on matrix size of 15000

```

Above: double precision matrix times.  
Below: single precision matrix time.  
Factor of ~50 faster.

```

>>> exec(open("time_matrix_mult_gpu.py").read())
got 144.9786211783412 gigaflops on matrix size of 100 wi
got 3960.0037765538946 gigaflops on matrix size of 300 v
got 18293.769490786173 gigaflops on matrix size of 1000
got 23998.77257796772 gigaflops on matrix size of 3000 v
got 26182.762495318155 gigaflops on matrix size of 5000
got 25521.487883626043 gigaflops on matrix size of 7000
got 25363.916185408034 gigaflops on matrix size of 10000
got 25545.671098791634 qigaflops on matrix size of 15000

```



# Laplace Kernel

```
import numpy as np; print("using numpy")
#import cupy as np; print("using cupy")
from cupyx.profiler import benchmark

def kernel(V):
    return V-0.25*(np.roll(V,1,0)+np.roll(V,-1,0)+np.roll(V,1,1)+np.roll(V,-1,1))

n=2048
V=np.zeros([n,n],dtype='float32')
x0=n//2
V[x0,x0]=1
stats=benchmark(kernel,(V,),n_repeat=100)
print("cpu/gpu times (msec): ",np.mean(stats.cpu_times)*1000,np.mean(stats.gpu_times)*1000)
```

```
[>>> exec(open("laplace_roll.py").read())
using cupy
cpu/gpu times (msec):  0.2452143561094999  0.4908240002393723
[>>> exec(open("laplace_roll.py").read())
using numpy
cpu/gpu times (msec):  13.675076654180884  13.708930616378783
```



# CUDA Kernel

- Under the hood, CUDA is “almost” c++
- Much more control if you use.
- Much, much harder to get right/debug in general. Lots of idiosyncrasies with hardware/threads getting in each other’s way.
- Often required for peak performance.

# Adding Arrays

```
__global__  
void add_vecs_simple(float *out, float *in1, float *in2, long n)  
{  
    /* Simple way to add vectors where each element gets one thread.  works great  
       for small arrays.*/  
    long idx=threadIdx.x+blockDim.x*blockIdx.x;  
    out[idx]=in1[idx]+in2[idx];  
}
```

```
extern "C" {  
void add(float *out, float *in1, float *in2, long n)  
{  
    long bs=256; //Set a block size for threads per block  
    long nblock=n/bs;  
    if ((nblock*bs)<n)  
        nblock++;  
    add_vecs_simple<<<nblock,bs>>>(out,in1,in2,n);  
    printf("err is currently %s\n",cudaGetErrorString(cudaGetLastError()));  
}
```

